

TensorFlow in Action (Part 1)

Chapter 3 : Keras and Data Retrieval in TensorFlow 2

This chapter introduces **Keras** as the high-level API in TensorFlow 2 for defining and training neural networks, and then focuses on how to build robust input pipelines for those models. It starts with a small image-classification example and uses it to compare three Keras model-building styles: the Sequential API, the Functional API, and the subclassing API. With Sequential, you stack layers in a straight line; with the Functional API, you define arbitrary DAG-style architectures (multiple inputs/outputs, skip connections); and with subclassing, you write custom `tf.keras.Model` classes for maximum flexibility.

Once the basic model-building patterns are clear, the chapter turns to **data ingestion**, which is crucial because deep learning models expect large, well-structured batches rather than ad-hoc Python loops. First, it uses the `tf.data` API: you create `tf.data.Dataset` objects from file paths and labels, then chain transformations like `map`, `shuffle`, `batch`, and `prefetch` to build efficient pipelines that can decode images, normalize them, and feed them to `model.fit` in a streaming fashion. The example with a flower image dataset demonstrates how you can assemble an end-to-end pipeline (folder structure → dataset → preprocessing → batched tensors) that trains a small ConvNet with just a few lines of glue code.

The chapter then shows two alternative data-loading tools. Keras **DataGenerators**, such as `ImageDataGenerator`, provide a simpler but less flexible way to load and optionally augment images directly from directories or from a pandas DataFrame using helpers like `flow_from_dataframe`, which can quickly produce batches of tensors and labels from a CSV mapping filenames to classes. Finally, the `tensorflow-datasets` (`tfds`) package is introduced as the easiest option when you use standard benchmark datasets; you can load data like CIFAR-10 or IMDB reviews with one call to `tfds.load`, getting pre-split, preformatted datasets ready to plug into `tf.data` and Keras. The trade-off is clear: `tf.data` offers maximum flexibility and performance, Keras generators are convenient for many image tasks, and `tfds` is perfect when your dataset is already supported.